

# **Lecture 46**

## **Multiversion Concurrency Control Techniques**

# Multiversion Concurrency Control Technique

- These protocols for concurrency control keep copies of the old values of a data item when the item is updated (written);
- They are known as **multiversion concurrency control** because several versions (values) of an item are kept by the system.
- When a transaction requests to read an item, the *appropriate* version is chosen to maintain the serializability of the currently executing schedule.
- One reason for keeping multiple versions is that some read operations that would be rejected in other techniques can still be accepted by reading an *older version* of the item to maintain serializability.
- When a transaction writes an item, it writes a *new version* and the old version(s) of the item is retained.
- Some multiversion concurrency control algorithms use the concept of view serializability rather than conflict serializability.

# Multiversion Concurrency Control Technique

- An obvious drawback of multiversion techniques is that more storage is needed to maintain multiple versions of the database items.
- In some cases, older versions can be kept in a temporary store.
- It is also possible that older versions may have to be maintained anyway—for example, for recovery purposes.
- Some database applications may require older versions to be kept to maintain a history of the changes of data item values.
- The extreme case is a *temporal database*, which keeps track of all changes and the times at which they occurred.
- In such cases, there is no additional storage penalty for multiversion techniques, since older versions are already maintained.

# Multiversion Timestamp Ordering Protocol

- In this method, several versions  $X_1, X_2, \dots, X_k$  of each data item  $X$  are maintained.
- For *each version*, the value of version  $X_i$  and the following two timestamps associated with version  $X_i$  are kept:
  1. **read\_TS( $X_i$ )** → The **read timestamp** of  $X_i$  is the largest of all the timestamps of transactions that have successfully read version  $X_i$ .
  2. **write\_TS( $X_i$ )** → The **write timestamp** of  $X_i$  is the timestamp of the transaction that wrote the value of version  $X_i$ .
- Whenever a transaction  $T$  is allowed to execute a `write_item( $X$ )` operation, a new version  $X_{k+1}$  of item  $X$  is created, with both the `write_TS( $X_{k+1}$ )` and the `read_TS( $X_{k+1}$ )` set to `TS( $T$ )`. Correspondingly, when a transaction  $T$  is allowed to read the value of version  $X_i$ , the value of `read_TS( $X_i$ )` is set to the larger of the current `read_TS( $X_i$ )` and `TS( $T$ )`.

# Multiversion Timestamp Ordering Protocol

- To ensure serializability, the following rules are used:
  1. If transaction  $T$  issues a `write_item(X)` operation, and version  $i$  of  $X$  has the highest  $\text{write\_TS}(X_i)$  of all versions of  $X$  that is also *less than or equal to*  $\text{TS}(T)$ , and  $\text{read\_TS}(X_i) > \text{TS}(T)$ , then abort and roll back transaction  $T$ ; otherwise, create a new version  $X_j$  of  $X$  with  $\text{read\_TS}(X_j) = \text{write\_TS}(X_j) = \text{TS}(T)$ .
  2. If transaction  $T$  issues a `read_item(X)` operation, find the version  $i$  of  $X$  that has the highest  $\text{write\_TS}(X_i)$  of all versions of  $X$  that is also *less than or equal to*  $\text{TS}(T)$ ; then return the value of  $X_i$  to transaction  $T$ , and set the value of  $\text{read\_TS}(X_i)$  to the larger of  $\text{TS}(T)$  and the current  $\text{read\_TS}(X_i)$ .

# Multiversion Timestamp Ordering Protocol

- As we can see in case 2, a  $\text{read\_item}(X)$  is *always successful*, since it finds the appropriate version  $X_i$  to read based on the  $\text{write\_TS}$  of the various existing versions of  $X$ .
- In case 1, however, transaction  $T$  may be aborted and rolled back. This happens if  $T$  attempts to write a version of  $X$  that should have been read by another transaction  $T'$  whose timestamp is  $\text{read\_TS}(X_i)$ ; however,  $T'$  has already read version  $X_i$ , which was written by the transaction with timestamp equal to  $\text{write\_TS}(X_i)$ . If this conflict occurs,  $T$  is rolled back; otherwise, a new version of  $X$ , written by transaction  $T$ , is created.
- Note that if  $T$  is rolled back, cascading rollback may occur. Hence, to ensure recoverability, a transaction  $T$  should not be allowed to commit until after all the transactions that have written some version that  $T$  has read have committed.

# Multiversion Two-Phase Locking Using Certify Locks

- In this multiple-mode locking scheme, there are *three locking modes* for an item— read, write, and *certify*—instead of just the two modes (read, write).
- Hence, the state of LOCK( $X$ ) for an item  $X$  can be one of read-locked, write-locked, certify-locked, or unlocked.
- In the standard locking scheme, with only read and write locks, a write lock is an exclusive lock. We can describe the relationship between read and write locks in the standard scheme by means of the **lock compatibility table** shown below.

|       | Read | Write |
|-------|------|-------|
| Read  | Yes  | No    |
| Write | No   | No    |

# Multiversion Two-Phase Locking Using Certify Locks

- An entry of *Yes* means that if a transaction  $T$  holds the type of lock specified in the column header on item  $X$  and if transaction  $T'$  requests the type of lock specified in the row header on the same item  $X$ , then  $T'$  *can obtain the lock* because the locking modes are compatible.
- On the other hand, an entry of *No* in the table indicates that the locks are not compatible, so  $T'$  *must wait* until  $T$  *releases* the lock.

|       | Read | Write |
|-------|------|-------|
| Read  | Yes  | No    |
| Write | No   | No    |



# Multiversion Two-Phase Locking Using Certify Locks

- The lock compatibility table for this scheme is shown in Figure

|         | Read | Write | Certify |
|---------|------|-------|---------|
| Read    | Yes  | Yes   | No      |
| Write   | Yes  | No    | No      |
| Certify | No   | No    | No      |

# Multiversion Two-Phase Locking Using Certify Locks

- In the standard locking scheme, once a transaction obtains a write lock on an item, no other transactions can access that item.
- The idea behind multiversion 2PL is to allow other transactions  $T'$  to read an item  $X$  while a single transaction  $T$  holds a write lock on  $X$ . This is accomplished by allowing *two versions* for each item  $X$ ; one version, the **committed version**, must always have been written by some committed transaction.
- The second **local version**  $X'$  can be created when a transaction  $T$  acquires a write lock on  $X$ . Other transactions can continue to read the *committed version* of  $X$  while  $T$  holds the write lock.

# Multiversion Two-Phase Locking Using Certify Locks

- Transaction  $T$  can write the value of  $X'$  as needed, without affecting the value of the committed version  $X$ . However, once  $T$  is ready to commit, it must obtain a **certify lock** on all items that it currently holds write locks on before it can commit; this is another form of **lock upgrading**.
- The certify lock is not compatible with read locks, so the transaction may have to delay its commit until all its write-locked items are released by any reading transactions in order to obtain the certify locks.
- Once the certify locks—which are exclusive locks—are acquired, the committed version  $X$  of the data item is set to the value of version  $X'$ , version  $X'$  is discarded, and the certify locks are then released.

For Video lecture on this topic please subscribe to my youtube channel.

The link for my youtube channel is

[https://www.youtube.com/channel/UCRWGtE76JITp1iim6aOTRuW?sub\\_confirmation=1](https://www.youtube.com/channel/UCRWGtE76JITp1iim6aOTRuW?sub_confirmation=1)